

Using fashion_mnist dataset and training the model with SimpleRNN.

```
import matplotlib.pyplot as plt
from tensorflow.keras import datasets, layers, models

(x_train, y_train), (x_test, y_test) = datasets.fashion_mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = models.Sequential([
    layers.SimpleRNN(64, input_shape=(28,28), activation='relu'),
    layers.Dense(10, activation='softmax')
])

model.summary()
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs = 3, validation_data = (x_test, y_test), batch_size = 64)
test_loss, test_acc = model.evaluate(x_train, y_train, batch_size = 64)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz> 32768/29515 [=====] - 0s 0us/step
40960/29515 [=====] - 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz> 26427392/26421880 [=====] - 0s 0us/step
26435584/26421880 [=====] - 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz> 16384/5148 [=====] - 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz> 4423680/4422102 [=====] - 0s 0us/step
4431872/4422102 [=====] - 0s 0us/step
Model: "sequential_1"

Layer (type)	Output Shape	Param #
simple_rnn_1 (SimpleRNN)	(None, 64)	5952
dense_1 (Dense)	(None, 10)	650

=====
Total params: 6,602
Trainable params: 6,602
Non-trainable params: 0

Epoch	1/3
938/938	[=====] - 9s 9ms/step - loss: 0.8554 - accuracy: 0.6805 - val_loss: 0.6252 - val_accuracy: 0.7804
Epoch 2/3	
938/938	[=====] - 8s 9ms/step - loss: 0.5704 - accuracy: 0.7990 - val_loss: 0.6056 - val_accuracy: 0.8002
Epoch 3/3	
938/938	[=====] - 8s 9ms/step - loss: 0.5054 - accuracy: 0.8218 - val_loss: 0.5479 - val_accuracy: 0.8019

Replace SimpleRNN with LSTM. The result is better than the preview's result.

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 64)	23808
dense_2 (Dense)	(None, 10)	650

=====
Total params: 24,458
Trainable params: 24,458
Non-trainable params: 0

Epoch	1/3
938/938	[=====] - 22s 22ms/step - loss: 0.7025 - accuracy: 0.7412 - val_loss: 0.4904 - val_accuracy: 0.8172
Epoch 2/3	
938/938	[=====] - 21s 22ms/step - loss: 0.4374 - accuracy: 0.8388 - val_loss: 0.4298 - val_accuracy: 0.8441
Epoch 3/3	
938/938	[=====] - 21s 22ms/step - loss: 0.3901 - accuracy: 0.8559 - val_loss: 0.4022 - val_accuracy: 0.8534
938/938	[=====] - 9s 9ms/step - loss: 0.3674 - accuracy: 0.8625

Add more LSTM layers.

✓
2m

```
import matplotlib.pyplot as plt
from tensorflow.keras import datasets, layers, models

(x_train, y_train), (x_test, y_test) = datasets.fashion_mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = models.Sequential([
    layers.LSTM(64, input_shape=(28,28), activation='relu', return_sequences = True),
    layers.LSTM(64, activation = 'relu'),
    layers.Dense(10, activation='softmax')
])

model.summary()
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs = 3, validation_data = (x_test, y_test), batch_size = 64)
test_loss, test_acc = model.evaluate(x_train, y_train, batch_size = 64)
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
lstm_1 (LSTM)	(None, 28, 64)	23808
lstm_2 (LSTM)	(None, 64)	33024
dense_3 (Dense)	(None, 10)	650

=====
Total params: 57,482
Trainable params: 57,482
Non-trainable params: 0

Epoch 1/3
938/938 [=====] - 44s 45ms/step - loss: 0.7053 - accuracy: 0.7400 - val_loss: 0.5401 - val_accuracy: 0.8059
Epoch 2/3
938/938 [=====] - 42s 45ms/step - loss: 0.4432 - accuracy: 0.8347 - val_loss: 0.4516 - val_accuracy: 0.8272
Epoch 3/3
938/938 [=====] - 42s 45ms/step - loss: 0.3867 - accuracy: 0.8548 - val_loss: 0.3927 - val_accuracy: 0.8559
938/938 [=====] - 15s 16ms/step - loss: 0.3565 - accuracy: 0.8673

Add extra Dense layers. It helps push accuracy a little bit more.

✓
2m

```
import matplotlib.pyplot as plt
from tensorflow.keras import datasets, layers, models

(x_train, y_train), (x_test, y_test) = datasets.fashion_mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = models.Sequential([
    layers.LSTM(64, input_shape=(28,28), activation='relu', return_sequences = True),
    layers.LSTM(64, activation = 'relu'),
    layers.Dense(32, activation = 'relu'),
    layers.Dense(10, activation='softmax')
])

model.summary()
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs = 3, validation_data = (x_test, y_test), batch_size = 64)
test_loss, test_acc = model.evaluate(x_train, y_train, batch_size = 64)
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
lstm_3 (LSTM)	(None, 28, 64)	23808
lstm_4 (LSTM)	(None, 64)	33024
dense_4 (Dense)	(None, 32)	2080
dense_5 (Dense)	(None, 10)	330

=====
Total params: 59,242
Trainable params: 59,242
Non-trainable params: 0
=====
Epoch 1/3
938/938 [=====] - 45s 46ms/step - loss: 0.7112 - accuracy: 0.7472 - val_loss: 0.4932 - val_accuracy: 0.8121
Epoch 2/3
938/938 [=====] - 41s 44ms/step - loss: 0.4314 - accuracy: 0.8410 - val_loss: 0.4300 - val_accuracy: 0.8462
Epoch 3/3
938/938 [=====] - 41s 44ms/step - loss: 0.3793 - accuracy: 0.8596 - val_loss: 0.3879 - val_accuracy: 0.8569
938/938 [=====] - 15s 16ms/step - loss: 0.3479 - accuracy: 0.8692

Add dropouts to avoid overfitting.

```
import matplotlib.pyplot as plt
from tensorflow.keras import datasets, layers, models

(x_train, y_train), (x_test, y_test) = datasets.fashion_mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = models.Sequential([
    layers.LSTM(64, input_shape=(28,28), activation='relu', return_sequences = True),
    layers.LSTM(64, activation = 'relu', dropout = 0.2, recurrent_dropout = 0.2),
    layers.Dense(32, activation = 'relu'),
    layers.Dropout(rate = 0.1),
    layers.Dense(10, activation='softmax')
])

model.summary()
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs = 3, validation_data = (x_test, y_test), batch_size = 64)
test_loss, test_acc = model.evaluate(x_test, y_test, batch_size = 64)
```

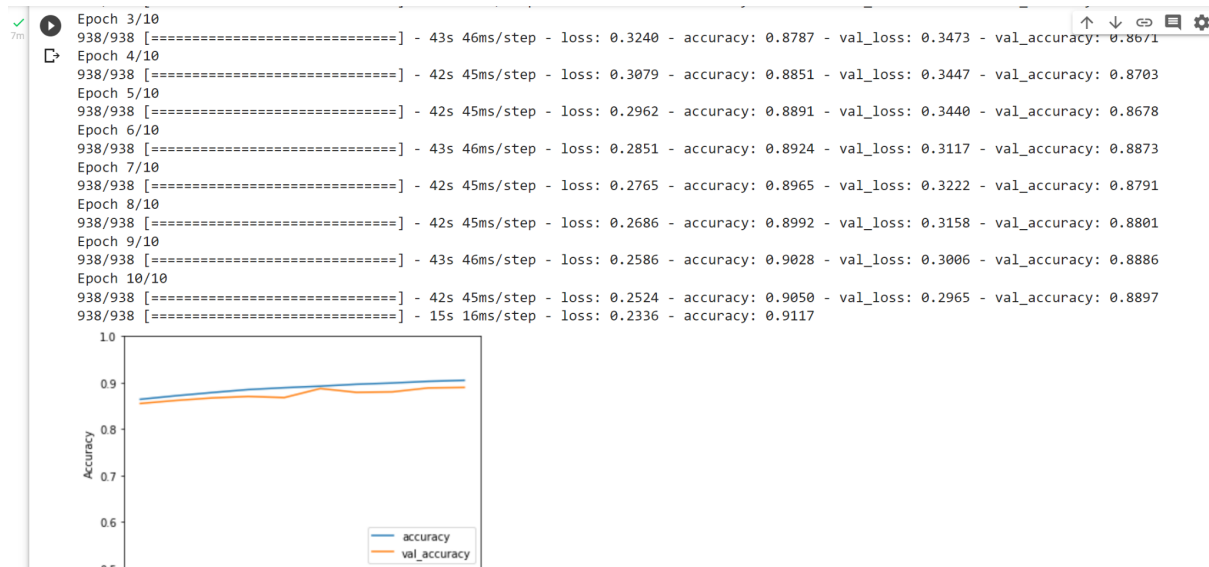
Model: "sequential_5"

Layer (type)	Output Shape	Param #
lstm_5 (LSTM)	(None, 28, 64)	23808
lstm_6 (LSTM)	(None, 64)	33024
dense_6 (Dense)	(None, 32)	2080
dropout (Dropout)	(None, 32)	0
dense_7 (Dense)	(None, 10)	330

=====
Total params: 59,242
Trainable params: 59,242
Non-trainable params: 0
=====
Epoch 1/3
938/938 [=====] - 74s 75ms/step - loss: 0.7887 - accuracy: 0.7030 - val_loss: 0.5164 - val_accuracy: 0.8058
Epoch 2/3
938/938 [=====] - 69s 74ms/step - loss: 0.4937 - accuracy: 0.8208 - val_loss: 0.4575 - val_accuracy: 0.8313
Epoch 3/3
938/938 [=====] - 70s 75ms/step - loss: 0.4441 - accuracy: 0.8392 - val_loss: 0.4269 - val_accuracy: 0.8455
938/938 [=====] - 14s 15ms/step - loss: 0.4023 - accuracy: 0.8554

Visualize the accuracy.

The result without adding dropout layers.



The result that added dropout layers.

